

Supercomputing: An Overview

Shelley Knuth

shelley.knuth@colorado.edu

Peter Ruprecht

peter.ruprecht@colorado.edu

www.rc.colorado.edu

Link to survey on this topic: <http://goo.gl/forms/8VidcwOhRT>

Slides: https://github.com/ResearchComputing/Final_Tutorials

What Is a Supercomputer?

- A supercomputer is one large computer made up of many smaller computers and processors
- Each different computer is called a node
- Each node has processors/cores
 - Carry out the instructions of the computer
- With a supercomputer, all these different computers talk to each other through a communications network
 - Example - InfiniBand

Computers and Cars - Analogy



Computers and Cars - Analogy



Why Use a Supercomputer?

- Supercomputers give you the opportunity to solve problems that are too complex for the desktop
 - Might take hours, days, weeks, months, years
 - If you use a supercomputer, might only take minutes, hours, days, or weeks
- Useful for problems that require large amounts of memory

World's Fastest Supercomputers

www.top500.org June 2015

Rank	Site	Name	TeraFlops
1	National Super Computer Center (Guangzhou, China)	Tianhe-2	54902.4
2	Oak Ridge National Laboratory (United States)	Titan	27112.5
3	DOE/NNSA/LLNL (United States)	Sequoia	20132.7
4	RIKEN Advanced Institute for Computational Science (Japan)	K	11280.4
5	DOE/Argonne National Lab (United States)	Mira	10066.3
6	Swiss National Supercomputing Centre (Switzerland)	Piz Daint	7788.9
7	King Abdullah University of Science and Technology (Saudi Arabia)	Shaheen II	7235.2
8	Texas Advanced Computing Center (United States)	Stampede	8520.1
9	Forschungszentrum Juelich (Germany)	JUQUEEN	5872.0
10	DOE/NNSA/LLNL (United States)	Vulcan	5033.2

What Does It Mean to Be Fast?

- Titan can do 27 trillion calculations per second
- A regular PC can perform 17 billion per second
- Researchers can get access to some of these systems through XSEDE (The Extreme Science and Engineering Discovery Environment)

Different Node Types

- Login nodes
 - This is where you are when you log in
 - No heavy computation, interactive jobs, or long running processes
 - Script or code editing, minor compiling
 - Job submission
- Compute/batch nodes
 - This is where jobs that are submitted through the scheduler run
 - Intended for heavy computation

Storage Spaces

- System variations
- **Home Directories**
 - Store source code
 - Not for direct computation
 - Small quota (~5 GB)
 - Backed up
- **\$WORK Space**
 - Mid level quota (~300 GB)
 - Large file storage
 - Not backed up

- **Scratch Directory**
 - Much larger – depends on system
 - Output from running jobs should go here
 - Files generally purged at some point

What is Job Scheduling

- Supercomputers usually consist of many nodes
- Users submit jobs that may run on one or multiple nodes
- Sometimes these jobs are very large; sometimes there are many small jobs
- Need software that will distribute the jobs appropriately
 - Make sure the job requirements are met
 - Reserve nodes until enough are available to run a job
 - Account for offline nodes
- Also need software to manage the resources
- Integrated with scheduler

Job Scheduling

- On a supercomputer, jobs are scheduled rather than just run instantly at the command line
 - People “buy” time to use the resources
 - Shared system
 - Request the amount of resources needed and for how long
 - Jobs are put in a queue until resources are available
 - Once the job is run they are “charged” for the time they used

Job Scheduling - Priority

- What jobs receive priority?
 - Can depend on the center
 - Can arrange for certain people who “pay more” receive priority
 - Generally though based on job size and time of entry
- Might have different queues based on different job needs
- Can receive priority on a job by creating a reservation

Job Schedulers - Slurm

- Jobs on supercomputers are managed and run by different software
- Simple Linux Utility for Resource Management (Slurm)
 - Open source software package
- Slurm is a resource manager
 - Keeps track of what nodes are busy/available, and what jobs are queued or running
- Slurm is a scheduler
 - Tells the resource manager when to run which job on the available resources

Running Jobs

- What is a “job”?
- Interactive jobs
 - Work interactively at the command line of a compute node
- Batch jobs
 - Submit job that will be executed when resources are available
 - Create a text file containing information about the job
 - Submit the job file to a queue
- Load the Slurm module!

Queues

- There are several ways to define a “queue”
- Clusters may have different queues set up to run different types of jobs
 - Certain queues might exist on certain clusters/resources
 - Other queues might be limited by maximum wall time
- Slurm can use a “quality of service” for each queue
 - aka “QOS”
- Also can use a “partition” (or set of nodes) that corresponds to a queue

Submit Batch Job example

```
#!/bin/bash
```

```
#SBATCH -N 2
```

```
#SBATCH --ntasks-per-node=12
```

```
#SBATCH --time=1:00:00
```

```
#SBATCH --job-name=SLURMDemo
```

```
#SBATCH --output=SLURMDemo.out
```

```
###SBATCH -A <account>
```

```
###SBATCH --mail-type=end
```

```
###SBATCH --mail-user=<your@email>
```

```
#No. nodes
```

```
#No. cores
```

```
#Max walltime
```

```
#Job name
```

```
#Output file name
```

```
#Allocation
```

```
#Send Email completion
```

```
#Email address
```

```
ml intel
```

```
ml openmpi/1.8.5
```

```
mpirun ./hello
```


Submit Batch Job example

- Have to make sure the slurm module is loaded!
- Submit the job, and specify the queue:
`sbatch --qos janus-debug slurmSub.sh`
 - Demonstrates that you can add slurm functions at the command line or in the bash script
- Check job status in the janus-debug queue:
`squeue -q janus-debug`
- Check output:
`cat SLURMDemo.out`

Your Turn

- Submit a slurm job with the following instructions:
 1. The job should run the Unix “hostname” command
 - Hint – the command “srun” will run commands in slurm
 2. The job should be submitted from a bash script named practice.sh
 - Don’t forget to make it executable!
 3. The job should run for 5 minutes in the default queue
 4. The job should be run on 1 node
 5. The output should be put in a file called hostname.txt

Your Turn - Solution

Bash Script practice.sh:

```
#!/bin/bash
#SBATCH -N 1                # No. of nodes
#SBATCH --time=0:05:00      # Walltime
#SBATCH --output=hostname.txt # Output file name

srun hostname
```

Submit the job:

```
sbatch practice.sh
```

Questions?

- Email rc-help@colorado.edu
- Twitter: CUBoulderRC
- Link to survey on this topic:
<http://goo.gl/forms/8VidcwOhRT>
- Slides:
https://github.com/ResearchComputing/Final_Tutorials