

Parallel Computing with Examples (MPI)

Shelley Knuth, Research Computing, University of Colorado-Boulder

shelley.knuth@colorado.edu

Questions? #RC_Meetups

Link to survey on this topic: <http://goo.gl/forms/8VidcwOhRT>


Slides: https://github.com/ResearchComputing/Final_Tutorials

Outline

- Distributed memory
- What is MPI?
- How is MPI used?
- Communicating
- Examples


Programming to Use Parallelism

- Parallelism across processors/threads
 - OpenMP
- Parallelism across multiple nodes - MPI



www.scan.co.uk

Distributed-memory Model



Distributed memory requires a communication network to connect memory

Programmers explicitly define how processors access other processor's memory

Source: https://computing.llnl.gov/tutorials/parallel_comp/#ModelsShared

MPI

- MPI is a library specification for message passing
- Widely used standard
- Can run on shared, distributed, or hybrid memory models
- Exchange data between processes through communication between tasks – send and receive data
- MPI can get complicated
- Programmers must explicitly implement parallelism using MPI constructs
- Portable

General MPI Code Structure

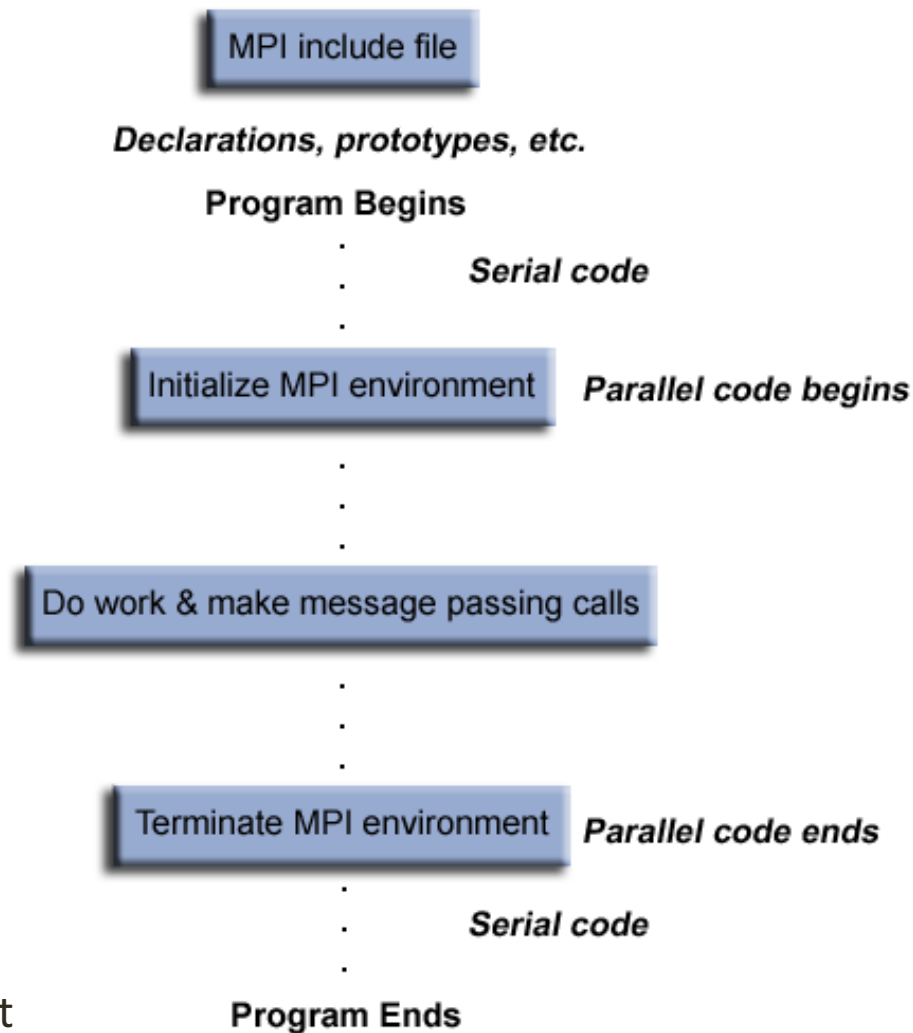
- You must have your header file at the top of any script you develop that uses MPI
- For C:

```
#include mpi.h
```

- For Fortran:

```
include mpif.h
```

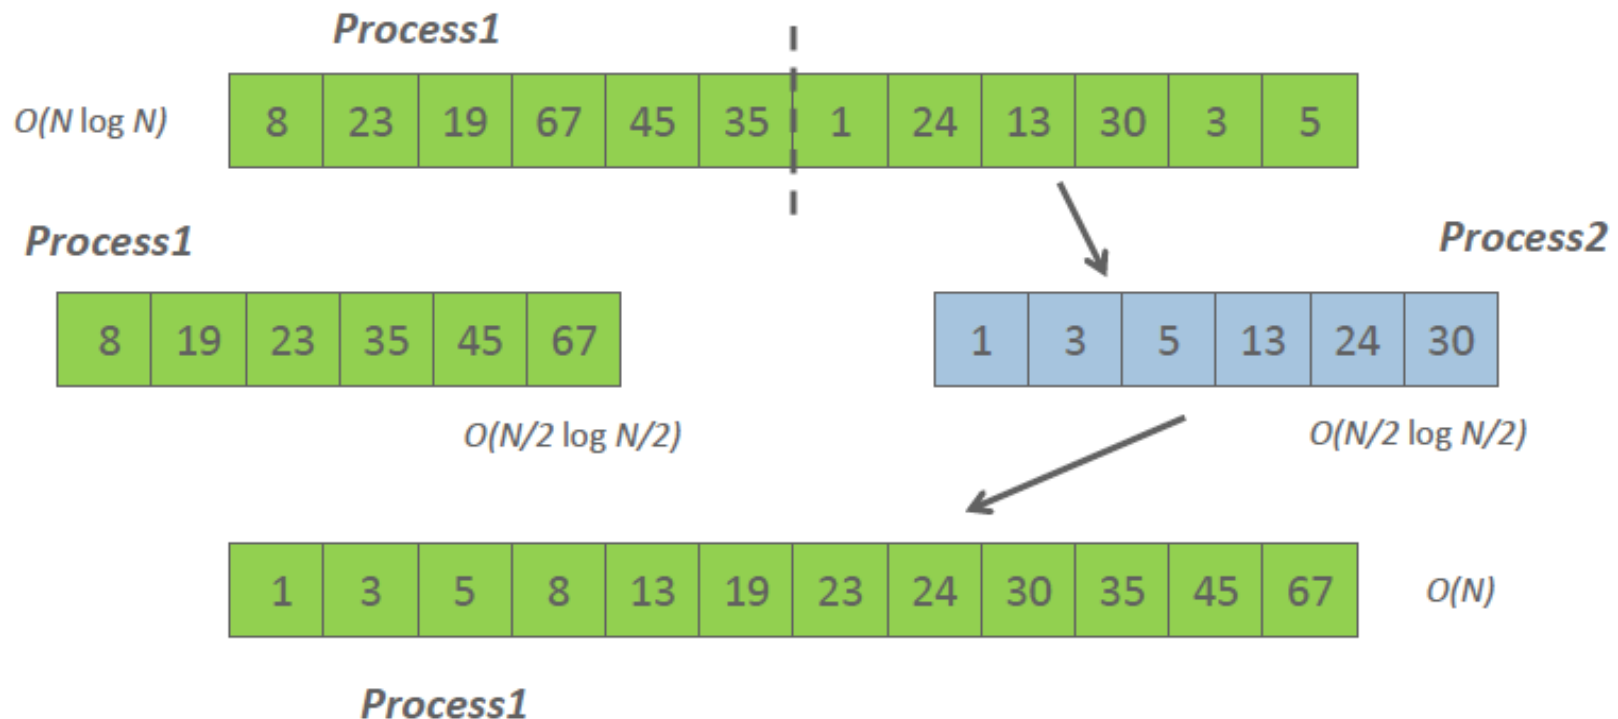
<https://computing.llnl.gov/tutorials/mpi/#What>



Message Passing

- A program that runs on a node is called a **process**
- When a program is run a process is run on each processor in the cluster
- These processes communicate with each other using message passing
- Message passing allows us to copy data from the memory of one process into another
- Message passing systems must at a minimum support system calls for sending and receiving messages


Example – Sorting Integers



http://hlor.inf.ethz.ch/teaching/mpl_tutorials/ppopp13/2013-02-24-ppopp-mpl-basic.pdf

MPI Communicators

- Communicators used to group collections of processes allowed to communicate with each other
- Assigns integers to each process at initialization
 - Called “rank”
- Programmer uses rank to specify destination or source for sending/receiving
- Initially all processes grouped into MPI_COMM_WORLD




<https://www.rc.usf.edu/tutorials/classes/tutorial/mpi/chapter2.html>

Environment Management Routines

- These routines set the MPI execution environment, and cover many purposes
- Some common routines:
 - MPI_INIT
 - MPI_COMM_SIZE
 - MPI_COMM_RANK
 - MPI_FINALIZE

How Do I Write A Program in MPI?

- Application needs to specify:
 - How do you compile and run the MPI application?
 - How will the processes be identified?
 - How will the data be described?



http://htr.inf.ethz.ch/teaching/mpi_tutorials/ppopp13/2013-02-24-ppopp-mpi-basic.pdf

Compiling and Running an MPI Application

- MPI applications can be written in C, C++, or Fortran and appropriate calls to MPI can be added where required

- Compiling code:

- Regular code:

- ```
gcc test.c -o test
```

- ```
ifort test.f -o test
```

- MPI applications:

- ```
mpicc test.c -o test
```

- ```
mpifort test.f -o test
```

- Running code:

- Regular code:

- ```
./test
```

- MPI applications (running with 16 processes):

- ```
mpiexec -np 16 ./test
```

MPI Library on Janus

- Unlike OpenMP, with MPI you need to have the appropriate library loaded in your environment
- Research Computing recommends impi
- To load these, just type:

```
ml gcc  
then  
ml impi
```

At the command line

Compiling An Application

- Before compiling an application, you MUST:
- Include the MPI header file
 - Needed to use all the MPI Library calls
- Initialize the MPI environment
 - `MPI_INIT()`
- Specify an end to the MPI environment at end of program
 - `MPI_Finalize()`

Example Fortran Code

Fortran code: simple.f90

To run:

```
ml slurm
ml gcc
ml impi
sinteractive --reservation=meetup
mpif90 simple.f90 -o simple
mpiexec -np 8 ./simple
```

OpenMP vs. MPI

Fortran code: hello.f90

The same code we ran as OpenMP modified for MPI

To run:

```
mpif90 hello.f90 -o hello  
mpiexec -np 8 ./hello
```


Communication

- One process sends a copy of data to another process and that process receives it
- Requires the following information
 - Sender needs to know
 - Who to send the data to
 - What kind of data to send
 - A tag (like an email subject) so the receiver understands what's being sent
 - Receiver maybe needs to know
 - Who is sending the data
 - What kind of data is sending
 - The tag

MPI_SEND (Fortran)

- MPI_SEND(buf, count, datatype, dest, tag, comm, ierr)
- Basic sending operation
- Routine returns only after the application buffer in the sending task is free for reuse
 - In some sense, a send cannot complete without acknowledgment from the receiving process
 - Can be changed
 - Out of scope here

What does this mean?

- **Buffer:** Usually variable name that is to be sent/received
- **Count:** number of data elements of a particular type to be sent
- **Datatype:** pre-defined data type of data (MPI_CHARACTER, MPI_INTEGER, etc)
- **Dest:** destination – indicates the process where the message should be delivered. Sent as the rank of the receiving process
- MPI_SEND(buf, count, datatype, dest, tag, comm, ierr)

What does this mean?

- **Tag:** Arbitrary number assigned by the programmer to identify a message.
 - **Comm:** communicator. Usually MPI_COMM_WORLD
 - **Ierr:** error message
-
- MPI_SEND(buf, count, datatype, dest, tag, comm, ierr)

MPI_RECV (Fortran)

- MPI_RECV(buf, count, datatype, source, tag, comm, status, ierr)
- **Status:** implies the source of the message
 - Integer array the size of MPI_STATUS_SIZE
- **Tag:** Can use MPI_ANY_TAG to receive any message regardless of tag

MPI Communication

Fortran code: ping.f90

To run:

```
mpif90 ping.f90 -o ping  
mpiexec -np 8 ./ping
```

References

Material for this talk is used from

- <https://computing.llnl.gov/tutorials/mpi/>
- http://htor.inf.ethz.ch/teaching/mpi_tutorials/ppopp13/2013-02-24-ppopp-mpi-basic.pdf
- <https://www.rc.usf.edu/tutorials/classes/tutorial/mpi/>
- These are great tutorials – we encourage you to go there for more information!

Questions?

- Email rc-help@colorado.edu
- Twitter: @CUBoulderRC
- Link to survey on this topic:
<http://goo.gl/forms/8VidcwOhRT>
- Slides:
https://github.com/ResearchComputing/Final_Tutorials
- Questions? #RC_Meetup