

# Parallel Computing with Examples (OpenMP)

Shelley Knuth, Research Computing, University of Colorado-Boulder

[shelley.knuth@colorado.edu](mailto:shelley.knuth@colorado.edu)

Questions? #RC\_Meetups

Link to survey on this topic: <http://goo.gl/forms/8VidcwOhRT>


Slides: [https://github.com/ResearchComputing/Final\\_Tutorials](https://github.com/ResearchComputing/Final_Tutorials)

# Outline

- Shared memory
- What is OpenMP?
- How is OpenMP used?
- Parallel region
- Public/Private variables
- Examples


# Programming to Use Parallelism

- Parallelism across processors/threads
  - OpenMP
- Parallelism across multiple nodes - MPI



[www.scan.co.uk](http://www.scan.co.uk)

# Shared-memory Model




The concept is that all processors can access all memory available

Multiple processors can perform tasks on their own but share the same memory

Source: [https://computing.llnl.gov/tutorials/parallel\\_comp/#ModelsShared](https://computing.llnl.gov/tutorials/parallel_comp/#ModelsShared)

# Shared-memory Model



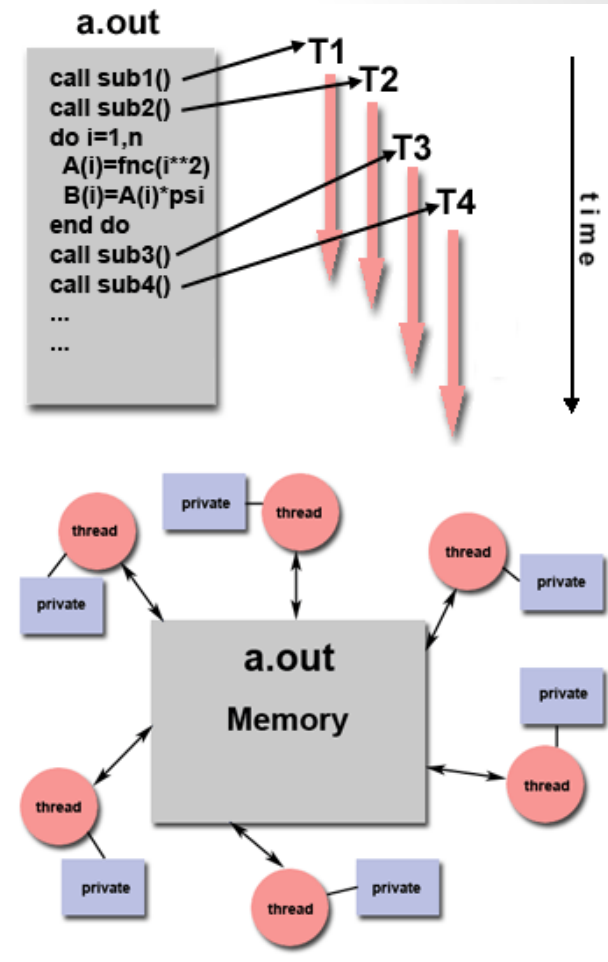
A thread is a block of code with one entry and one exit that is abstract and is mapped onto a physical core. Multiple threads can be mapped onto one core.

Threads communicate by depositing contents in shared memory area

Source: [http://people.math.umass.edu/~johnston/PHI\\_WG\\_2014/OpenMPslides\\_tamu\\_sc.pdf](http://people.math.umass.edu/~johnston/PHI_WG_2014/OpenMPslides_tamu_sc.pdf)

# Multi-Threaded, Shared Memory Parallelism

- Main program does many things, including run subroutines
  - Threads that can be run concurrently
  - Share the same resources from the main program, but also has local data
  - Threads communicate through global memory
  - Must ensure multiple threads don't update concurrently
  - Where OpenMP and programmers come in




Source: [https://computing.llnl.gov/tutorials/parallel\\_comp/#ModelsShared](https://computing.llnl.gov/tutorials/parallel_comp/#ModelsShared)

# OpenMP

- OpenMP: An application programming interface (API) for parallel programming on multiprocessors
- Uses shared memory
- OpenMP is used through compiler directives embedded in Fortran, C, or C++ code
- Directs multi-threaded, shared memory parallelism
- Can do a lot with only a handful of commands
- Intended to be easy to use

# OpenMP – Fork/Join

- OpenMP programs start with a single thread (master)
- Then Master creates a team of parallel “worker” threads (FORK)
- Statements in block are executed in parallel by every thread
- At end, all threads synchronize and join master thread





# OpenMP Directives

- Comments in source code that specify parallelism for shared memory machines
  - Enclosing parallel directives
- **FORTRAN:** directives begin with **!\$OMP**, **C\$OMP** or **\*\$OMP**
- **C/C++:** directives begin with **#pragma omp**

# OpenMP Fortran: General Code Structure – Parallel Regions

Parallel regions are blocks of code that will be executed by multiple threads

```
1      !$OMP PARALLEL
2          code block
3          call work(...)
4      !$OMP END PARALLEL
```

Line 1	Team of threads formed at parallel region.
Lines 2-3	Each thread executes code block and subroutine calls. No branching (in or out) in a parallel region.
Line 4	All threads synchronize at end of parallel region (implied barrier).


Use the thread number to divide work among threads.

# Parallel Regions

- When thread hits PARALLEL directive, creates team of threads
  - Becomes master
  - Code is duplicated and all threads execute that code
  - Runs the same code on different data
    - Split up loops and operate on different data
  - Only master thread continues after implied barrier
- Can determine number of threads by:
  - Setting the number threads to a default number or within code
  - Allowing number of threads to change from one parallel region to another


# Shared and Private Variables

- When specifying the **PRIVATE** clause, that variable is private to each thread
  - Each thread will have its own unique copy
  - Can only be accessed by the threads that own it
- If there are  $x$  team members, there are  $x+1$  copies of the variables in the private clause
  - One global copy and a private copy for each team member
- When specifying **SHARED** clause, all threads can access that data



# Shared and Private Variables

- Global variables are shared by default
- Variables declared within subroutines called within private region are by default private
  - Index variables are also default private
- Can change specifications with **DEFAULT** clause



# Shared and Private Variables

```
#pragma omp parallel for shared(a,b,c,n) private(i)
    for (i=0; i<n; i++){
        a[i] = b[i] + c[i];
    }
```

- All threads can access a, b, c, and n
- Each loop has own private copy of index i

# Private Variable Example

```
#pragma omp parallel for shared(a,b,c,n) private(temp,i)
    for (i=0; i< n; i++){
        temp = a[i] / b[i];
        c[i] = temp + cos(temp);
    }
```


- Variable temp also needs to be private
- Otherwise each thread would be reading/writing to same location

# Parallel Region Example

- Finding the integral
  - Area under a curve
  - Sum of the area of all the rectangles underneath the curve (approximate)

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$


$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$





# Parallel Region Example

- The same code is used to calculate the area of each of the rectangles
- Different threads will calculate different rectangles
- Which rectangles are calculated with each thread is random



# OpenMP Compiling

- When compiling must use appropriate compiler flag to turn on OpenMP compilations

Compiler / Platform	Compiler	Flag
Intel Linux Opteron/Xeon	icc icpc ifort	-openmp
PGI Linux Opteron/Xeon	pgcc pgCC pgf77 pgf90	-mp
GNU Linux Opteron/Xeon IBM Blue Gene	gcc g++ g77 gfortran	-fopenmp
IBM Blue Gene	bgxlc_r, bgcc_r bgxlC_r, bgxlc++_r bgxlc89_r bgxlc99_r bgxlf_r bgxlf90_r bgxlf95_r bgxlf2003_r *Be sure to use a thread-safe compiler - its name ends with _r	-qsmp=omp

# Runtime Library Routines

Routine	Purpose
<u>OMP SET NUM THREADS</u>	Sets the number of threads that will be used in the next parallel region
<u>OMP GET NUM THREADS</u>	Returns the number of threads that are currently in the team executing the parallel region from which it is called
<u>OMP GET THREAD NUM</u>	Returns the thread number of the thread, within the team, making this call.
<u>OMP GET THREAD LIMIT</u>	Returns the maximum number of OpenMP threads available to a program

- In Fortran some routines are functions; some are subroutines
- In C/C++, all are subroutines. Must include the omp.h header file

Fortran	<b>INTEGER FUNCTION OMP_GET_NUM_THREADS()</b>
C/C++	<b>#include &lt;omp.h&gt;</b> <b>int omp_get_num_threads(void)</b>

# OMP Code Practice – Exercise 1

- Code:

```
omp_hello.f
```

- Instructions for running:

```
ssh tutorial-login.rc.colorado.edu -l user00XX  
ml gcc/5.1.0  
ml slurm  
sinteractive --reservation=meetup  
gfortran -fopenmp omp_hello.f -o hello  
./hello
```

# How Do I Prepare My Code for OpenMP?

- I have code! I want it to be parallel too!
- Steps to go through
  1. Verify that code is parallelizable
    - Make sure you don't have any loop dependencies
  2. Analyze your code
    - Where does the program spend most of its time?
    - Look for loops
      - Typically easy to parallelize
      - Outside of nested loops

# How Do I Prepare My Code for OpenMP?

- Steps to go through

## 3. Restructure code

- Put `parallel do` constructs around parallelizable loops
- List variables with appropriate `shared`, `private`, etc. clauses
- Many other things you can do that we don't cover here

## 4. Overhead

- How much time was spent preparing your code for parallelization?
- Is this more than the time spent running your code serially?

# Example Code – Exercise 2

- Code:  
for.c
- Instructions for running:

```
ssh tutorial-login.rc.colorado.edu -l user00XX  
ml gcc/5.1.0  
ml slurm  
sinteractive --reservation=meetup
```

```
gcc for.c -o for_noflag  
time ./for_noflag 10000000
```

```
gcc -fopenmp for.c -o for  
time ./for 10000000
```

# Example Code – Exercise 2

- We need to consider whether our code really does experience a speed up
  - Array size 10,000,000
  - Drops by ~30%
- Also, what are we looking at for overhead times?
  - Array size 10
  - Takes longer for parallel code to run



# References

- [https://portal.tacc.utexas.edu/c/document\\_library/get\\_file?uuid=c3c38847-ca7e-41bf-aefa-fb232a777699&groupId=13601](https://portal.tacc.utexas.edu/c/document_library/get_file?uuid=c3c38847-ca7e-41bf-aefa-fb232a777699&groupId=13601)
- <https://computing.llnl.gov/tutorials/openMP/#Introduction>
- <http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>
- <https://computing.llnl.gov/tutorials/mpi/>

# Questions?

- Email [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Twitter: @CUBoulderRC
- Link to survey on this topic:  
<http://goo.gl/forms/8VidcwOhRT>
- Slides:  
[https://github.com/ResearchComputing/Final\\_Tutorials](https://github.com/ResearchComputing/Final_Tutorials)
- Questions? #RC\_Meetup