

Janus Supercomputer Bootcamp – Updates and Some Parallel Computing

Shelley Knuth and Pete Ruprecht, Research Computing, University of Colorado-Boulder

shelley.knuth@colorado.edu

peter.ruprecht@colorado.edu

Link to survey on this topic: <http://goo.gl/forms/8VidcwOhRT>

Slides: https://github.com/ResearchComputing/Final_Tutorials

Outline

- Update on awarded Janus MRI Proposal
- UCD BiPM HPC Update
- Parallel Programming Examples
- One-On-One Interactions and Tutorials

Next-Generation Supercomputer at CU-Boulder

- Funded via an NSF MRI grant awarded jointly to CU-Boulder and CSU
- \$2M to CU and \$700K to CSU ... with matching funds the hardware budget is about \$3.5M
- RFP has been published, vendor award by end of November (hopefully!)
- Installed and running late spring 2016

Next-Generation Supercomputer

- Expected performance about 450 TFLOPS (compared to about 170 for Janus)
- Compute nodes
 - Expect 24 real cores and 128 GB RAM
- 10 GPU/visualization nodes
 - 2x NVIDIA K80 GPUs
- 5 High-memory nodes
- 20 Xeon Phi (“Knight’s Landing”) nodes
- “Omni-Path” high-performance interconnect
- 1 PB of high-performance scratch storage

Next-Generation Supercomputer

- 10% of CPU-hours available to RMACC
- Opportunities to buy additional nodes to improve your job priority and guarantee an allocation
- “Name the supercomputer” contest
 - Best suggestion wins a Kindle Fire tablet
 - Send entries to becky.yeager@colorado.edu by 10/31

UCD BiPM HPC Update


Parallel Computing with Examples (OpenMP)

Outline

- Shared memory
- What is OpenMP?
- How is OpenMP used?
- Parallel region
- Public/Private variables
- Examples

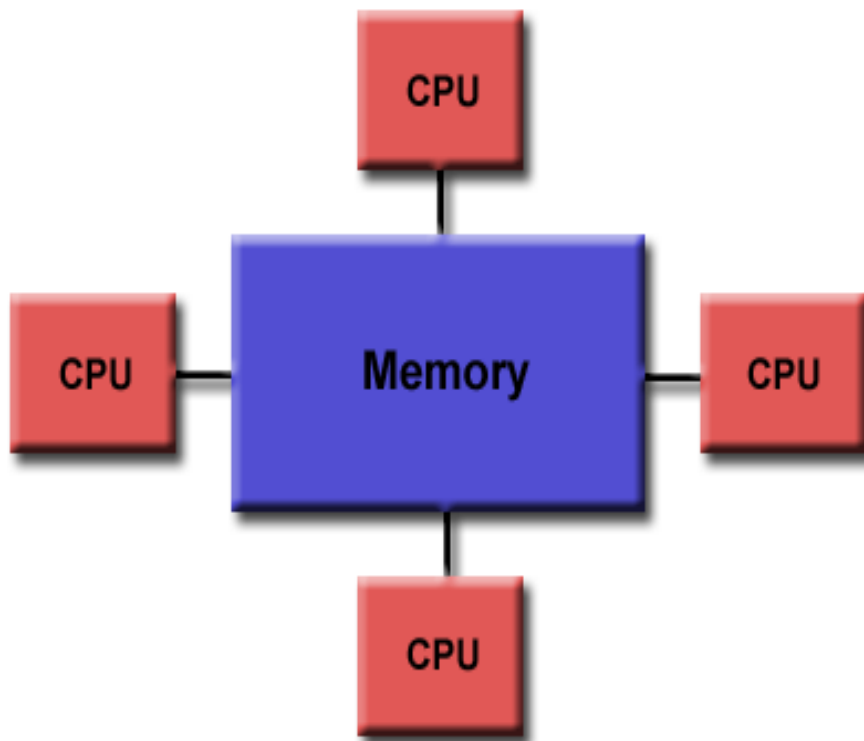
Programming to Use Parallelism

- Parallelism across processors/threads
 - OpenMP
- Parallelism across multiple nodes -
MPI



www.scan.co.uk

Shared-memory Model



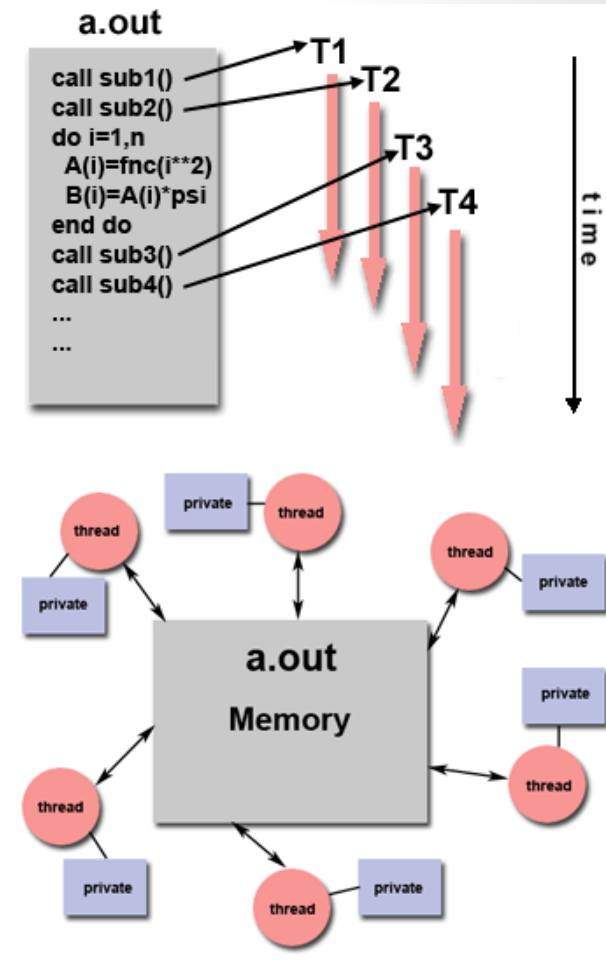
The concept is that all processors can access all memory available

Multiple processors can perform tasks on their own but share the same memory

Source: https://computing.llnl.gov/tutorials/parallel_comp/#ModelsShared

Multi-Threaded, Shared Memory Parallelism

- Main program does many things, including run subroutines
 - Threads that can be run concurrently
 - Share the same resources from the main program, but also has local data
 - Threads communicate through global memory
 - Must ensure multiple threads don't update concurrently
 - Where OpenMP and programmers come in




Source: https://computing.llnl.gov/tutorials/parallel_comp/#ModelsShared

OpenMP

- OpenMP: An application programming interface (API) for parallel programming on multiprocessors
- Uses shared memory
- OpenMP is used through compiler directives embedded in Fortran, C, or C++ code
- Directs multi-threaded, shared memory parallelism
- Can do a lot with only a handful of commands
- Intended to be easy to use

OpenMP – Fork/Join

- OpenMP programs start with a single thread (master)
- Then Master creates a team of parallel “worker” threads (FORK)
- Statements in block are executed in parallel by every thread
- At end, all threads synchronize and join master thread



OpenMP Directives

- Comments in source code that specify parallelism for shared memory machines
 - Enclosing parallel directives
- **FORTRAN:** directives begin with **!\$OMP**, **C\$OMP** or ***\$OMP**
- **C/C++:** directives begin with **#pragma omp**

OpenMP Fortran: General Code Structure – Parallel Regions

Parallel regions are blocks of code that will be executed by multiple threads

```
1      !$OMP PARALLEL
2          code block
3          call work(...)
4      !$OMP END PARALLEL
```

Line 1	Team of threads formed at parallel region.
Lines 2-3	Each thread executes code block and subroutine calls. No branching (in or out) in a parallel region.
Line 4	All threads synchronize at end of parallel region (implied barrier).


Use the thread number to divide work among threads.

Parallel Regions

- When thread hits PARALLEL directive, creates team of threads
 - Becomes master
 - Code is duplicated and all threads execute that code
 - Runs the same code on different data
 - Split up loops and operate on different data
 - Only master thread continues after implied barrier
- Can determine number of threads by:
 - Setting the number threads to a default number or within code
 - Allowing number of threads to change from one parallel region to another

Shared and Private Variables

- When specifying the **PRIVATE** clause, that variable is private to each thread
 - Each thread has own unique copy
 - Can only be accessed by the threads that own it
 - Variables declared in private subroutines are default private
 - Index variables are also default private
- When specifying **SHARED** clause, all threads can access that data
 - Global variables are shared by default



Private Variable Example

```
#pragma omp parallel for shared(a,b,c,n) private(temp,i)
    for (i=0; i< n; i++){
        temp = a[i] / b[i];
        c[i] = temp + cos(temp);
    }
```


- All threads can access a, b, c, and n
- Each loop has own private copy of index i
- Variable temp also needs to be private
- Otherwise each thread would be reading/writing to same location

Parallel Region Example

- Finding the integral
 - Area under a curve
 - Sum of the area of all the rectangles underneath the curve (approximate)


$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$



Parallel Region Example

- The same code is used to calculate the area of each of the rectangles
- Different threads will calculate different rectangles
- Which rectangles are calculated with each thread is random



OpenMP Compiling

- When compiling must use appropriate compiler flag to turn on OpenMP compilations

Compiler / Platform	Compiler	Flag
Intel Linux Opteron/Xeon	icc icpc ifort	-openmp
PGI Linux Opteron/Xeon	pgcc pgCC pgf77 pgf90	-mp
GNU Linux Opteron/Xeon IBM Blue Gene	gcc g++ g77 gfortran	-fopenmp
IBM Blue Gene	bgxlc_r, bgcc_r bgxlC_r, bgxlc++_r bgxlc89_r bgxlc99_r bgxlf_r bgxlf90_r bgxlf95_r bgxlf2003_r *Be sure to use a thread-safe compiler - its name ends with _r	-qsmp=omp

Runtime Library Routines

Routine	Purpose
<u>OMP SET NUM THREADS</u>	Sets the number of threads that will be used in the next parallel region
<u>OMP GET NUM THREADS</u>	Returns the number of threads that are currently in the team executing the parallel region from which it is called
<u>OMP GET THREAD NUM</u>	Returns the thread number of the thread, within the team, making this call.
<u>OMP GET THREAD LIMIT</u>	Returns the maximum number of OpenMP threads available to a program

- In C/C++, must include the omp.h header file

Fortran	INTEGER FUNCTION OMP_GET_NUM_THREADS()
C/C++	#include <omp.h> int omp_get_num_threads(void)

OMP Code Practice – Exercise 1

- Code:

```
omp_hello.f
```

- Instructions for running:

```
ssh tutorial-login.rc.colorado.edu -l user00XX  
ml intel  
ml slurm  
sinteractive --reservation=meetup  
ifort -qopenmp omp_hello.f -o hello  
./hello
```

How Do I Prepare My Code for OpenMP?

- I have code! I want it to be parallel too!
- Steps to go through
 1. Verify that code is parallelizable
 - Make sure you don't have any loop dependencies
 2. Analyze your code
 - Where does the program spend most of its time?
 - Look for loops
 - Typically easy to parallelize
 - Outside of nested loops

How Do I Prepare My Code for OpenMP?

- Steps to go through

3. Restructure code

- Put `parallel do` constructs around parallelizable loops
- List variables with appropriate `shared`, `private`, etc. clauses
- Many other things you can do that we don't cover here

4. Overhead

- How much time was spent preparing your code for parallelization?
- Is this more than the time spent running your code serially?

Example Code – Exercise 2

- Code:
for.c
- Instructions for running:

```
ssh tutorial-login.rc.colorado.edu -l user00XX  
ml intel  
ml slurm  
sinteractive --reservation=meetup
```

```
icc for.c -o for_noflag  
time ./for_noflag 10000000
```

```
icc -qopenmp for.c -o for  
time ./for 10000000
```

Example Code – Exercise 2


- We need to consider whether our code really does experience a speed up
 - Array size 10,000,000
 - Drops by ~30%
- Also, what are we looking at for overhead times?
 - Array size 10
 - Takes longer for parallel code to run

Parallel Computing with Examples (MPI)

Outline

- Distributed memory
- What is MPI?
- How is MPI used?
- Communicating
- Examples

Distributed-memory Model



Distributed memory requires a communication network to connect memory

Programmers explicitly define how processors access other processor's memory

Source: https://computing.llnl.gov/tutorials/parallel_comp/#ModelsShared

MPI

- MPI is a library specification for message passing
- Widely used standard
- Can run on shared, distributed, or hybrid memory models
- Exchange data between processes through communication between tasks – send and receive data
- MPI can get complicated
- Programmers must explicitly implement parallelism using MPI constructs
- Portable

General MPI Code Structure

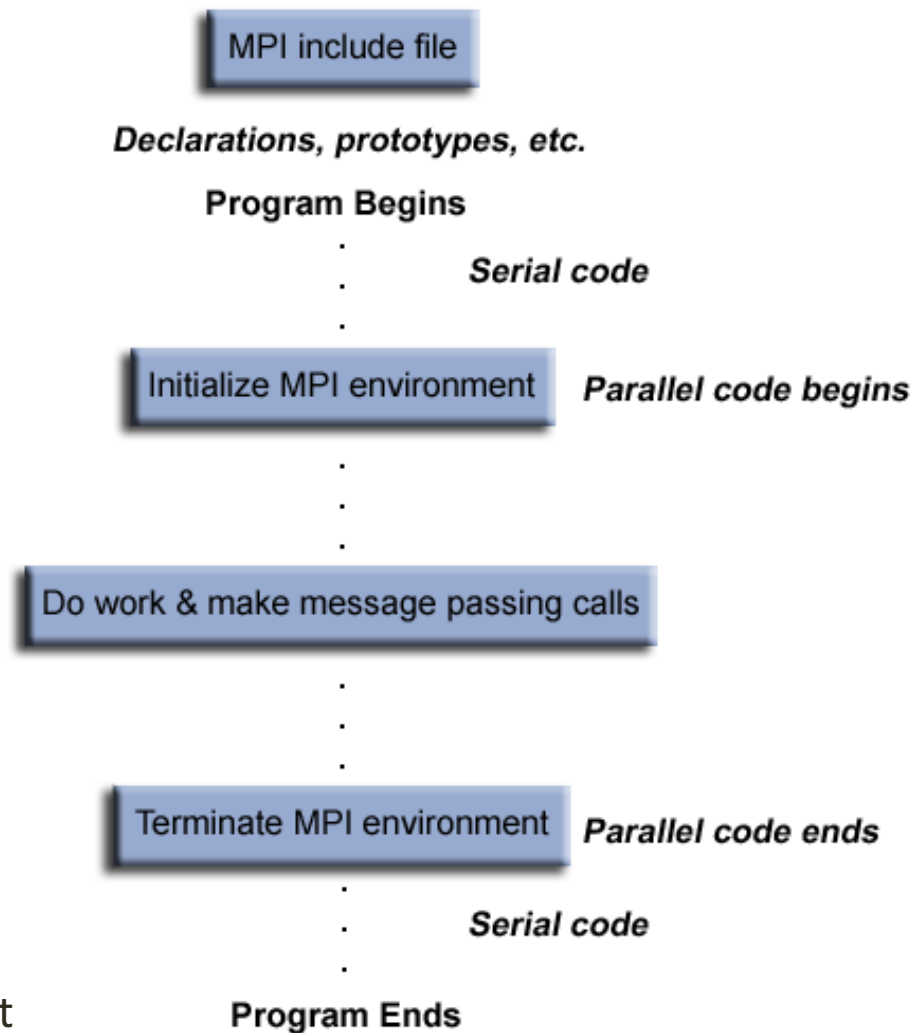
- You must have your header file at the top of any script you develop that uses MPI
- For C:

```
#include mpi.h
```

- For Fortran:

```
include mpif.h
```


<https://computing.llnl.gov/tutorials/mpi/#What>



Message Passing

- A program that runs on a node is called a **process**
- When a program is run a process is run on each processor in the cluster
- These processes communicate with each other using message passing
- Message passing allows us to copy data from the memory of one process into another
- Message passing systems must at a minimum support system calls for sending and receiving messages


Example – Sorting Integers



http://hlor.inf.ethz.ch/teaching/mpl_tutorials/ppopp13/2013-02-24-ppopp-mpl-basic.pdf

MPI Communicators

- Communicators used to group collections of processes allowed to communicate with each other
- Assigns integers to each process at initialization
 - Called “rank”
- Programmer uses rank to specify destination or source for sending/receiving
- Initially all processes grouped into MPI_COMM_WORLD




<https://www.rc.usf.edu/tutorials/classes/tutorial/mpi/chapter2.html>

Environment Management Routines

- These routines set the MPI execution environment, and cover many purposes
- Some common routines:
 - MPI_INIT
 - MPI_COMM_SIZE
 - MPI_COMM_RANK
 - MPI_FINALIZE

How Do I Write A Program in MPI?

- Application needs to specify:
 - How do you compile and run the MPI application?
 - How will the processes be identified?
 - How will the data be described?



http://hlor.inf.ethz.ch/teaching/mpl_tutorials/ppopp13/2013-02-24-ppopp-mpl-basic.pdf

Compiling and Running an MPI Application

- MPI applications can be written in C, C++, or Fortran and appropriate calls to MPI can be added where required

- Compiling code:

- Regular code:

```
gcc test.c -o test
```

```
ifort test.f -o test
```

- MPI applications:

```
mpicc test.c -o test
```

```
mpifort test.f -o test
```

- Running code:

- Regular code:

```
./test
```

- MPI applications (running with 16 processes):

```
mpiexec -np 16 ./test
```

MPI Library on Janus

- Unlike OpenMP, with MPI you need to have the appropriate library loaded in your environment
- Research Computing recommends impi
- To load these, just type:

```
ml gcc  
then  
ml impi
```

At the command line

Compiling An Application

- Before compiling an application, you MUST:
- Include the MPI header file
 - Needed to use all the MPI Library calls
- Initialize the MPI environment
 - `MPI_INIT()`
- Specify an end to the MPI environment at end of program
 - `MPI_Finalize()`

Example Fortran Code

Fortran code: simple.f90

To run:

```
ml slurm
ml gcc
ml impi
sinteractive --reservation=meetup
mpif90 simple.f90 -o simple
mpiexec -np 8 ./simple
```

OpenMP vs. MPI

Fortran code: hello.f90

The same code we ran as OpenMP modified for MPI

To run:

```
mpif90 hello.f90 -o hello  
mpiexec -np 8 ./hello
```

Time?

Communication

- One process sends a copy of data to another process and that process receives it
- Requires the following information
 - Sender needs to know
 - Who to send the data to
 - What kind of data to send
 - A tag (like an email subject) so the receiver understands what's being sent
 - Receiver maybe needs to know
 - Who is sending the data
 - What kind of data is sending
 - The tag

MPI_SEND (Fortran)

- MPI_SEND(buf, count, datatype, dest, tag, comm, ierr)
- Basic sending operation
- Routine returns only after the application buffer in the sending task is free for reuse
 - In some sense, a send cannot complete without acknowledgment from the receiving process
 - Can be changed
 - Out of scope here

What does this mean?

- **Buffer:** Usually variable name that is to be sent/received
- **Count:** number of data elements of a particular type to be sent
- **Datatype:** pre-defined data type of data (MPI_CHARACTER, MPI_INTEGER, etc)
- **Dest:** destination – indicates the process where the message should be delivered. Sent as the rank of the receiving process
- MPI_SEND(buf, count, datatype, dest, tag, comm, ierr)

What does this mean?

- **Tag:** Arbitrary number assigned by the programmer to identify a message.
 - **Comm:** communicator. Usually MPI_COMM_WORLD
 - **Ierr:** error message
-
- MPI_SEND(buf, count, datatype, dest, tag, comm, ierr)

MPI_RECV (Fortran)

- MPI_RECV(buf, count, datatype, source, tag, comm, status, ierr)
- **Status:** implies the source of the message
 - Integer array the size of MPI_STATUS_SIZE
- **Tag:** Can use MPI_ANY_TAG to receive any message regardless of tag

MPI Communication

Fortran code: ping.f90

To run:

```
mpif90 ping.f90 -o ping  
mpiexec -np 8 ./ping
```

References

- https://portal.tacc.utexas.edu/c/document_library/get_file?uuid=c3c38847-ca7e-41bf-aefa-fb232a777699&groupId=13601
- <https://computing.llnl.gov/tutorials/openMP/>
- <http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>
- <https://computing.llnl.gov/tutorials/mpi/>
- http://htor.inf.ethz.ch/teaching/mpi_tutorials/ppopp13/2013-02-24-ppopp-mpi-basic.pdf
- <https://www.rc.usf.edu/tutorials/classes/tutorial/mpi/>

Questions?

- Email rc-help@colorado.edu
- Twitter: @CUBoulderRC
- Link to survey on this topic:
<http://goo.gl/forms/8VidcwOhRT>
- Slides:
https://github.com/ResearchComputing/Final_Tutorials
- Questions? #RC_Meetup